

KP02-Copy1

November 14, 2021

```
[1]: !pip install pyspark
```

```
Requirement already satisfied: pyspark in /usr/lib/spark/python (3.1.2)
Requirement already satisfied: py4j==0.10.9 in
/opt/conda/miniconda3/lib/python3.8/site-packages (from pyspark) (0.10.9)
WARNING: Running pip as the 'root' user can result in broken permissions
and conflicting behaviour with the system package manager. It is recommended to
use a virtual environment instead: https://pip.pypa.io/warnings/venv
```

```
[2]: import os
import pandas as pd
import numpy as np

from pyspark import SparkConf, SparkContext
from pyspark.sql import SparkSession, SQLContext

import pyspark.sql.functions as F
from pyspark.sql.functions import udf, col

from pyspark.ml.tuning import ParamGridBuilder, CrossValidator
from pyspark.ml.evaluation import BinaryClassificationEvaluator

from pyspark.sql.functions import udf, col

from pyspark.mllib.evaluation import RegressionMetrics

from pyspark.ml.tuning import ParamGridBuilder, CrossValidator,
↳CrossValidatorModel
from pyspark.ml.feature import VectorAssembler, StandardScaler
from pyspark.ml.evaluation import RegressionEvaluator

from pyspark.ml.tuning import ParamGridBuilder, CrossValidator
from pyspark.ml.evaluation import BinaryClassificationEvaluator

from pyspark.ml import Pipeline
```

```
[3]: from pyspark.sql import SparkSession
      spark = SparkSession.builder.getOrCreate()
```

```
[4]: df = spark.read.format('csv').option("header", "true").option("delimiter", ",").
      ↪load("gs://bdamlproject/WA_Fn-UseC_-HR-Employee-Attrition.csv")
```

```
[5]: for col in df.columns:
      print(col)
```

Age
Attrition
BusinessTravel
DailyRate
Department
DistanceFromHome
Education
EducationField
EmployeeCount
EmployeeNumber
EnvironmentSatisfaction
Gender
HourlyRate
JobInvolvement
JobLevel
JobRole
JobSatisfaction
MaritalStatus
MonthlyIncome
MonthlyRate
NumCompaniesWorked
Over18
OverTime
PercentSalaryHike
PerformanceRating
RelationshipSatisfaction
StandardHours
StockOptionLevel
TotalWorkingYears
TrainingTimesLastYear
WorkLifeBalance
YearsAtCompany
YearsInCurrentRole
YearsSinceLastPromotion
YearsWithCurrManager

```
[6]: df.count()
```

```
[6]: 1470
```

```
[7]: df.dtypes
```

```
[7]: [('Age', 'string'),
      ('Attrition', 'string'),
      ('BusinessTravel', 'string'),
      ('DailyRate', 'string'),
      ('Department', 'string'),
      ('DistanceFromHome', 'string'),
      ('Education', 'string'),
      ('EducationField', 'string'),
      ('EmployeeCount', 'string'),
      ('EmployeeNumber', 'string'),
      ('EnvironmentSatisfaction', 'string'),
      ('Gender', 'string'),
      ('HourlyRate', 'string'),
      ('JobInvolvement', 'string'),
      ('JobLevel', 'string'),
      ('JobRole', 'string'),
      ('JobSatisfaction', 'string'),
      ('MaritalStatus', 'string'),
      ('MonthlyIncome', 'string'),
      ('MonthlyRate', 'string'),
      ('NumCompaniesWorked', 'string'),
      ('Over18', 'string'),
      ('OverTime', 'string'),
      ('PercentSalaryHike', 'string'),
      ('PerformanceRating', 'string'),
      ('RelationshipSatisfaction', 'string'),
      ('StandardHours', 'string'),
      ('StockOptionLevel', 'string'),
      ('TotalWorkingYears', 'string'),
      ('TrainingTimesLastYear', 'string'),
      ('WorkLifeBalance', 'string'),
      ('YearsAtCompany', 'string'),
      ('YearsInCurrentRole', 'string'),
      ('YearsSinceLastPromotion', 'string'),
      ('YearsWithCurrManager', 'string')]
```

```
[8]: df1=df.toPandas()
```

21/11/13 01:39:17 WARN org.apache.spark.sql.catalyst.util.package: Truncated the string representation of a plan since it was too large. This behavior can be

adjusted by setting 'spark.sql.debug.maxToStringFields'.

[9]: df1

```
[9]:      Age Attrition      BusinessTravel DailyRate      Department \
0      41      Yes      Travel_Rarely      1102      Sales
1      49      No      Travel_Frequently      279      Research & Development
2      37      Yes      Travel_Rarely      1373      Research & Development
3      33      No      Travel_Frequently      1392      Research & Development
4      27      No      Travel_Rarely      591      Research & Development
... ..      ...
1465   36      No      Travel_Frequently      884      Research & Development
1466   39      No      Travel_Rarely      613      Research & Development
1467   27      No      Travel_Rarely      155      Research & Development
1468   49      No      Travel_Frequently      1023      Sales
1469   34      No      Travel_Rarely      628      Research & Development

      DistanceFromHome Education EducationField EmployeeCount EmployeeNumber \
0                      1          2 Life Sciences              1              1
1                      8          1 Life Sciences              1              2
2                      2          2      Other              1              4
3                      3          4 Life Sciences              1              5
4                      2          1      Medical              1              7
... ..      ...
1465                  23      2      Medical              1            2061
1466                   6      1      Medical              1            2062
1467                   4      3 Life Sciences              1            2064
1468                   2      3      Medical              1            2065
1469                   8      3      Medical              1            2068

      ... RelationshipSatisfaction StandardHours StockOptionLevel \
0      ...                      1              80              0
1      ...                      4              80              1
2      ...                      2              80              0
3      ...                      3              80              0
4      ...                      4              80              1
... ..      ...
1465   ...                      3              80              1
1466   ...                      1              80              1
1467   ...                      2              80              1
1468   ...                      4              80              0
1469   ...                      1              80              0

      TotalWorkingYears TrainingTimesLastYear WorkLifeBalance YearsAtCompany \
0                      8              0              1              6
1                      10             3              3             10
2                      7              3              3              0
```

3	8	3	3	8
4	6	3	3	2
...
1465	17	3	3	5
1466	9	5	3	7
1467	6	0	3	6
1468	17	3	2	9
1469	6	3	4	4

	YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager
0	4	0	5
1	7	1	7
2	0	0	0
3	7	3	0
4	2	2	2
...
1465	2	0	3
1466	7	1	7
1467	2	0	3
1468	6	0	8
1469	3	1	2

[1470 rows x 35 columns]

```
[10]: df1.describe()
```

```
[10]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	\
count	1470	1470	1470	1470	1470	
unique	43	2	3	886	3	
top	35	No	Travel_Rarely	691	Research & Development	
freq	78	1233	1043	6	961	

	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	\
count	1470	1470	1470	1470	1470	
unique	29	5	6	1	1470	
top	2	3	Life Sciences	1	190	
freq	211	572	606	1470	1	

	RelationshipSatisfaction	StandardHours	StockOptionLevel	\
count	...	1470	1470	1470
unique	...	4	1	4
top	...	3	80	0
freq	...	459	1470	631

	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	YearsAtCompany	\
count	1470	1470	1470	1470	
unique	40	7	4	37	

top	10	2	3	5
freq	202	547	893	196

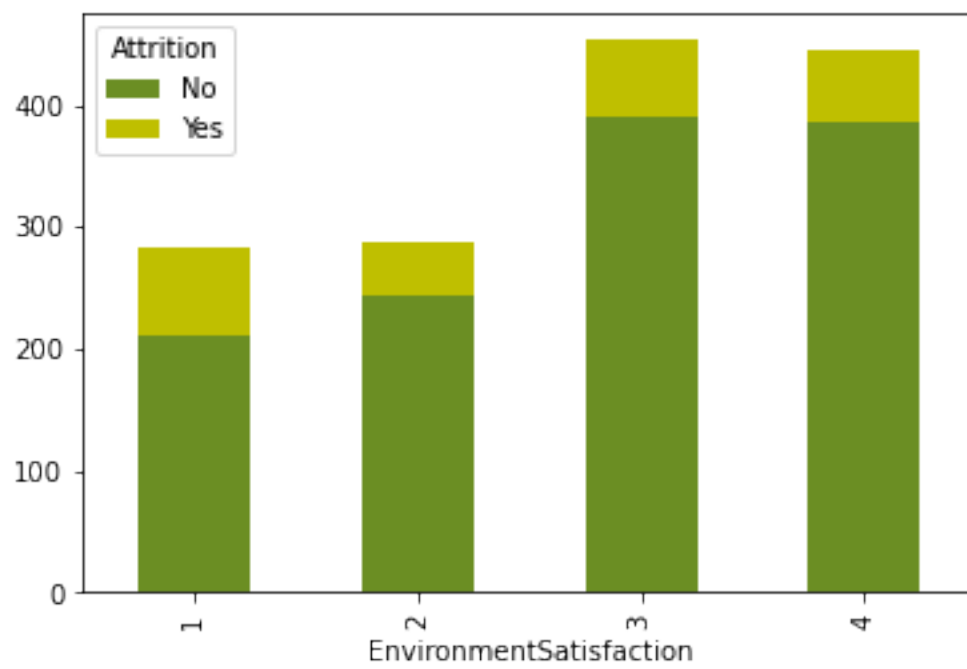
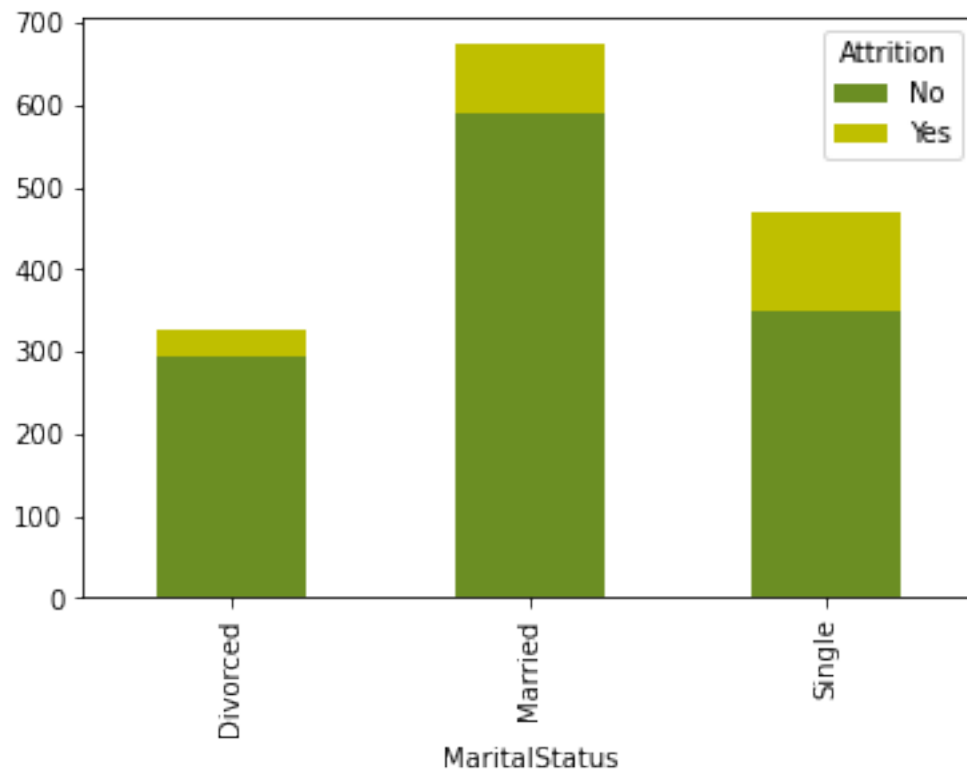
	YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager
count	1470	1470	1470
unique	19	16	18
top	2	0	2
freq	372	581	344

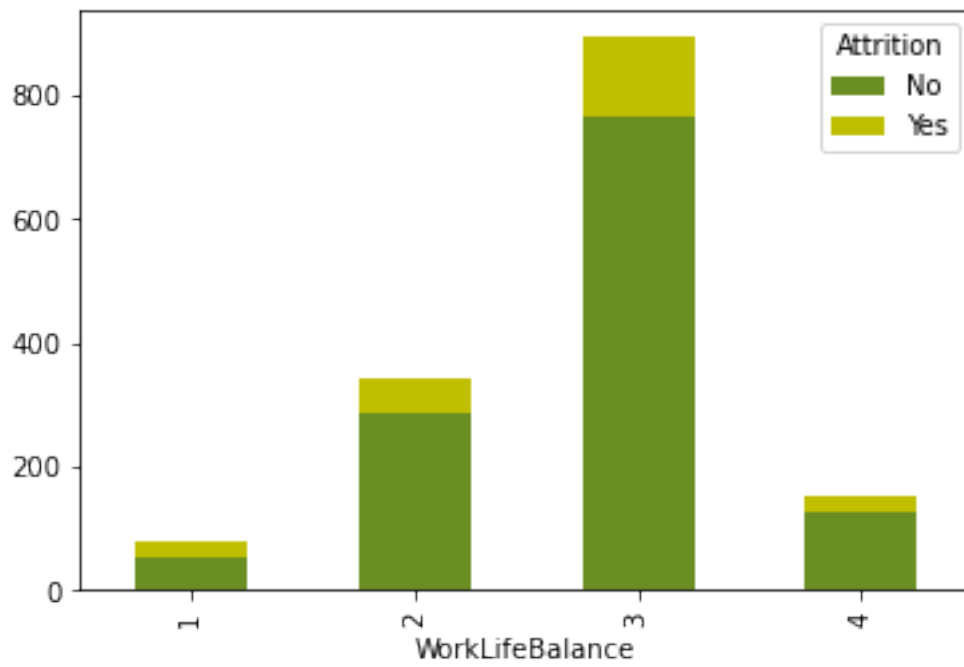
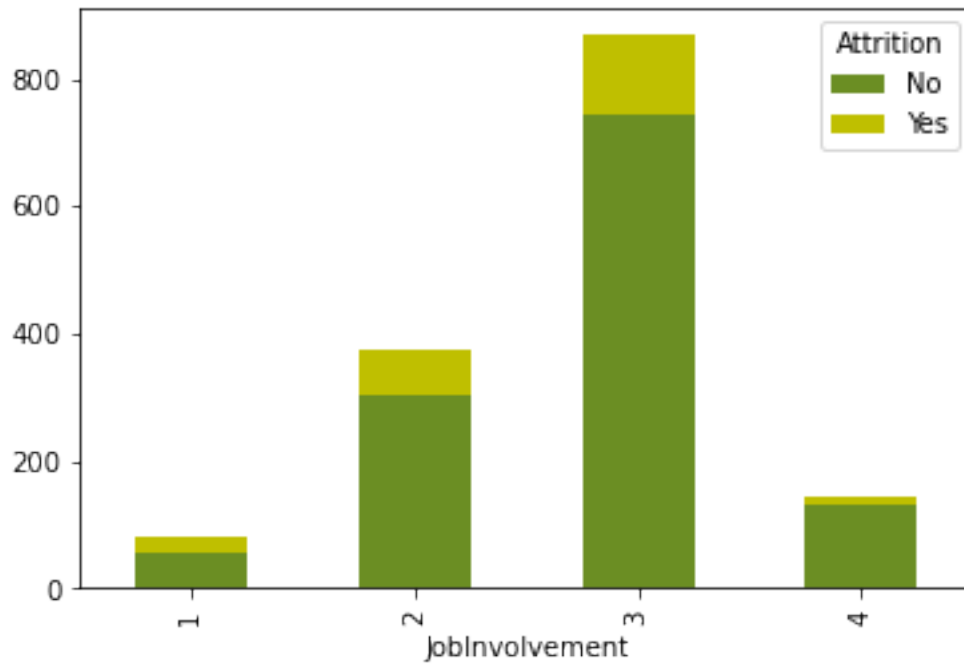
[4 rows x 35 columns]

```
[11]: df_plot_Mar = df1.groupby(['Attrition', 'MaritalStatus']).size().reset_index().
      ↪pivot(columns = 'Attrition', index = 'MaritalStatus', values = 0)
df_plot_Env = df1.groupby(['Attrition', 'EnvironmentSatisfaction']).size().
      ↪reset_index().pivot(columns = 'Attrition', index = 'EnvironmentSatisfaction', values = 0)
df_plot_Inv = df1.groupby(['Attrition', 'JobInvolvement']).size().reset_index().
      ↪pivot(columns = 'Attrition', index = 'JobInvolvement', values = 0)
df_plot_Bln = df1.groupby(['Attrition', 'WorkLifeBalance']).size().
      ↪reset_index().pivot(columns = 'Attrition', index = 'WorkLifeBalance', values = 0)
```

```
[12]: df_plot_Mar.plot(kind = 'bar', stacked = True, color = ['#6B8E23', 'y'])
df_plot_Env.plot(kind = 'bar', stacked = True, color = ['#6B8E23', 'y'])
df_plot_Inv.plot(kind = 'bar', stacked = True, color = ['#6B8E23', 'y'])
df_plot_Bln.plot(kind = 'bar', stacked = True, color = ['#6B8E23', 'y'])
```

```
[12]: <AxesSubplot:xlabel='WorkLifeBalance'>
```





```
[13]: df1.isna().sum()
```



```
[13]: Age 0
      Attrition 0
      BusinessTravel 0
      DailyRate 0
      Department 0
      DistanceFromHome 0
      Education 0
      EducationField 0
      EmployeeCount 0
      EmployeeNumber 0
      EnvironmentSatisfaction 0
      Gender 0
      HourlyRate 0
      JobInvolvement 0
      JobLevel 0
      JobRole 0
      JobSatisfaction 0
      MaritalStatus 0
      MonthlyIncome 0
      MonthlyRate 0
      NumCompaniesWorked 0
      Over18 0
      OverTime 0
      PercentSalaryHike 0
      PerformanceRating 0
      RelationshipSatisfaction 0
      StandardHours 0
      StockOptionLevel 0
      TotalWorkingYears 0
      TrainingTimesLastYear 0
      WorkLifeBalance 0
      YearsAtCompany 0
      YearsInCurrentRole 0
      YearsSinceLastPromotion 0
      YearsWithCurrManager 0
      dtype: int64
```

```
[14]: dfs = df1.groupby(['Attrition', 'Gender']).size().reset_index()
      dfs
```

```
[14]: Attrition Gender 0
      0 No Female 501
      1 No Male 732
      2 Yes Female 87
      3 Yes Male 150
```

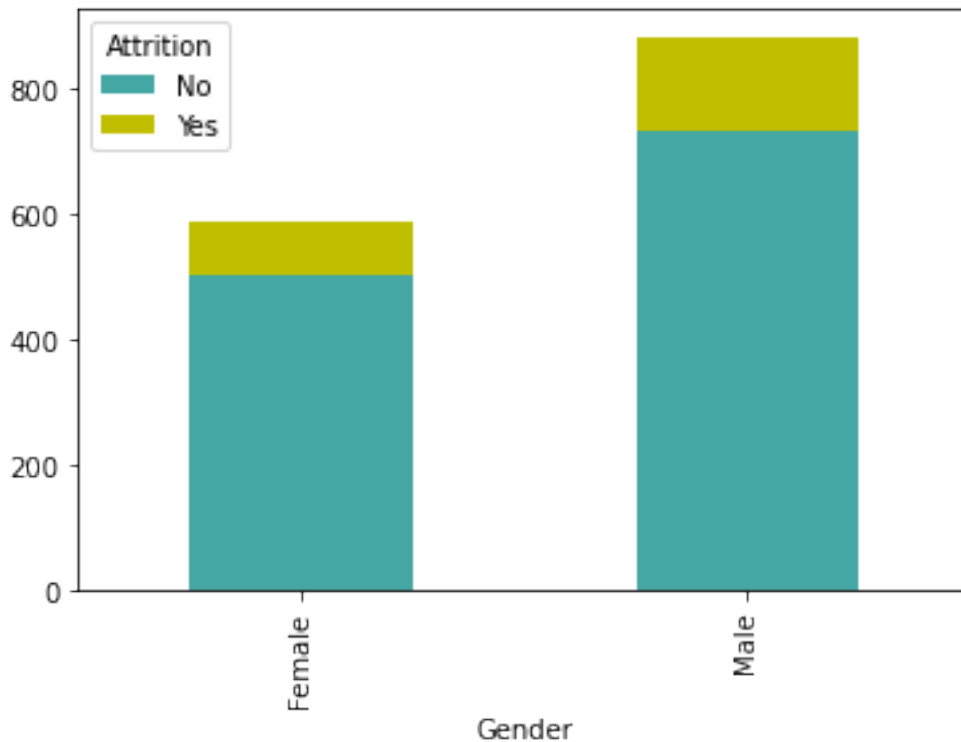
```
[15]: df_plot_Gen = df1.groupby(['Attrition', 'Gender']).size().reset_index().
      ↪pivot(columns = 'Attrition', index = 'Gender', values = 0)
df_plot_Gen
```

```
[15]: Attrition    No  Yes
      Gender
      Female    501   87
      Male      732  150
```

```
[16]: df_plot_Gen.plot(kind = 'bar', stacked = True, color = ['#45A7A3','y'])
plt.show()
```

```
-----
NameError                                Traceback (most recent call last)
/tmp/ipykernel_8047/759299616.py in <module>
      1 df_plot_Gen.plot(kind = 'bar', stacked = True, color = ['#45A7A3','y'])
----> 2 plt.show()

NameError: name 'plt' is not defined
```



```
[17]: from pyspark.ml.stat import Correlation
      from pyspark.ml.feature import VectorAssembler
```

```
[18]: #deleting columns that are not required
df_drop1=df.
↳drop('DailyRate','Department','BusinessTravel','Education','EmployeeCount','EmployeeNumber')
```

```
[19]: df_drop1.columns
```

```
[19]: ['Age',
      'Attrition',
      'DistanceFromHome',
      'EnvironmentSatisfaction',
      'Gender',
      'JobInvolvement',
      'JobLevel',
      'JobSatisfaction',
      'MaritalStatus',
      'MonthlyIncome',
      'NumCompaniesWorked',
      'PercentSalaryHike',
      'PerformanceRating',
      'TotalWorkingYears',
      'TrainingTimesLastYear',
      'WorkLifeBalance',
      'YearsAtCompany']
```

```
[20]: df_drop1.dtypes
```

```
[20]: [('Age', 'string'),
      ('Attrition', 'string'),
      ('DistanceFromHome', 'string'),
      ('EnvironmentSatisfaction', 'string'),
      ('Gender', 'string'),
      ('JobInvolvement', 'string'),
      ('JobLevel', 'string'),
      ('JobSatisfaction', 'string'),
      ('MaritalStatus', 'string'),
      ('MonthlyIncome', 'string'),
      ('NumCompaniesWorked', 'string'),
      ('PercentSalaryHike', 'string'),
      ('PerformanceRating', 'string'),
      ('TotalWorkingYears', 'string'),
      ('TrainingTimesLastYear', 'string'),
      ('WorkLifeBalance', 'string'),
      ('YearsAtCompany', 'string')]
```

```
[21]: from pyspark.sql.types import IntegerType
```

```
[22]: df_drop1= df_drop1.withColumn("Age",df["Age"].cast(IntegerType()))
df_drop1= df_drop1.withColumn("DistanceFromHome",df["DistanceFromHome"].
    ↳cast(IntegerType()))
df_drop1= df_drop1.
    ↳withColumn("EnvironmentSatisfaction",df["EnvironmentSatisfaction"].
    ↳cast(IntegerType()))
df_drop1= df_drop1.withColumn("JobInvolvement",df["JobInvolvement"].
    ↳cast(IntegerType()))
df_drop1= df_drop1.withColumn("JobLevel",df["JobLevel"].cast(IntegerType()))
df_drop1= df_drop1.withColumn("JobSatisfaction",df["JobSatisfaction"].
    ↳cast(IntegerType()))
df_drop1= df_drop1.withColumn("MonthlyIncome",df["MonthlyIncome"].
    ↳cast(IntegerType()))
df_drop1= df_drop1.withColumn("NumCompaniesWorked",df["NumCompaniesWorked"].
    ↳cast(IntegerType()))
df_drop1= df_drop1.withColumn("PercentSalaryHike",df["PercentSalaryHike"].
    ↳cast(IntegerType()))
df_drop1= df_drop1.withColumn("PerformanceRating",df["PerformanceRating"].
    ↳cast(IntegerType()))
df_drop1= df_drop1.withColumn("TotalWorkingYears",df["TotalWorkingYears"].
    ↳cast(IntegerType()))
df_drop1= df_drop1.
    ↳withColumn("TrainingTimesLastYear",df["TrainingTimesLastYear"].
    ↳cast(IntegerType()))
df_drop1= df_drop1.withColumn("WorkLifeBalance",df["WorkLifeBalance"].
    ↳cast(IntegerType()))
df_drop1= df_drop1.withColumn("YearsAtCompany",df["YearsAtCompany"].
    ↳cast(IntegerType()))
```

```
[23]: from pyspark.sql.functions import when
from pyspark.sql.functions import lit
```

```
[24]: df_drop1=df_drop1.withColumn("AttritionMod",when((df_drop1.Attrition=='Yes'),
    ↳lit(1)) .otherwise(lit(0)))
```

```
[25]: df_drop1.describe()
```

```
[25]: DataFrame[summary: string, Age: string, Attrition: string, DistanceFromHome:
string, EnvironmentSatisfaction: string, Gender: string, JobInvolvement: string,
JobLevel: string, JobSatisfaction: string, MaritalStatus: string, MonthlyIncome:
string, NumCompaniesWorked: string, PercentSalaryHike: string,
PerformanceRating: string, TotalWorkingYears: string, TrainingTimesLastYear:
string, WorkLifeBalance: string, YearsAtCompany: string, AttritionMod: string]
```

```
[26]: df_drop1.head(5)
```

```
[26]: [Row(Age=41, Attrition='Yes', DistanceFromHome=1, EnvironmentSatisfaction=2,
Gender='Female', JobInvolvement=3, JobLevel=2, JobSatisfaction=4,
MaritalStatus='Single', MonthlyIncome=5993, NumCompaniesWorked=8,
PercentSalaryHike=11, PerformanceRating=3, TotalWorkingYears=8,
TrainingTimesLastYear=0, WorkLifeBalance=1, YearsAtCompany=6, AttritionMod=1),
Row(Age=49, Attrition='No', DistanceFromHome=8, EnvironmentSatisfaction=3,
Gender='Male', JobInvolvement=2, JobLevel=2, JobSatisfaction=2,
MaritalStatus='Married', MonthlyIncome=5130, NumCompaniesWorked=1,
PercentSalaryHike=23, PerformanceRating=4, TotalWorkingYears=10,
TrainingTimesLastYear=3, WorkLifeBalance=3, YearsAtCompany=10, AttritionMod=0),
Row(Age=37, Attrition='Yes', DistanceFromHome=2, EnvironmentSatisfaction=4,
Gender='Male', JobInvolvement=2, JobLevel=1, JobSatisfaction=3,
MaritalStatus='Single', MonthlyIncome=2090, NumCompaniesWorked=6,
PercentSalaryHike=15, PerformanceRating=3, TotalWorkingYears=7,
TrainingTimesLastYear=3, WorkLifeBalance=3, YearsAtCompany=0, AttritionMod=1),
Row(Age=33, Attrition='No', DistanceFromHome=3, EnvironmentSatisfaction=4,
Gender='Female', JobInvolvement=3, JobLevel=1, JobSatisfaction=3,
MaritalStatus='Married', MonthlyIncome=2909, NumCompaniesWorked=1,
PercentSalaryHike=11, PerformanceRating=3, TotalWorkingYears=8,
TrainingTimesLastYear=3, WorkLifeBalance=3, YearsAtCompany=8, AttritionMod=0),
Row(Age=27, Attrition='No', DistanceFromHome=2, EnvironmentSatisfaction=1,
Gender='Male', JobInvolvement=3, JobLevel=1, JobSatisfaction=2,
MaritalStatus='Married', MonthlyIncome=3468, NumCompaniesWorked=9,
PercentSalaryHike=12, PerformanceRating=3, TotalWorkingYears=6,
TrainingTimesLastYear=3, WorkLifeBalance=3, YearsAtCompany=2, AttritionMod=0)]
```

```
[27]: df_drop1=df_drop1.withColumn("MaritalStatusMod",when((df_drop1.
↪MaritalStatus=='Yes'), lit(1)) .otherwise(lit(0)))
```

```
[28]: df_drop1=df_drop1.withColumn("GenderMod",when((df_drop1.Gender=='Male'),
↪lit(1)) .otherwise(lit(0)))
```

```
[29]: df_final=df_drop1.drop('Gender','MaritalStatus','Attrition')
```

```
[30]: df_final.columns
```

```
[30]: ['Age',
'DistanceFromHome',
'EnvironmentSatisfaction',
'JobInvolvement',
'JobLevel',
'JobSatisfaction',
'MonthlyIncome',
'NumCompaniesWorked',
'PercentSalaryHike',
```

```
'PerformanceRating',
'TotalWorkingYears',
'TrainingTimesLastYear',
'WorkLifeBalance',
'YearsAtCompany',
'AttritionMod',
'MaritalStatusMod',
'GenderMod']
```

```
[31]: df_final = df_final.withColumn("GenderMod",df_final["GenderMod"].
      ↳cast(IntegerType()))
df_final = df_final.withColumn("MaritalStatusMod",df_final["MaritalStatusMod"].
      ↳cast(IntegerType()))
df_final = df_final.withColumn("AttritionMod",df_final["AttritionMod"].
      ↳cast(IntegerType()))
```

```
[32]: df_final.dtypes
```

```
[32]: [('Age', 'int'),
      ('DistanceFromHome', 'int'),
      ('EnvironmentSatisfaction', 'int'),
      ('JobInvolvement', 'int'),
      ('JobLevel', 'int'),
      ('JobSatisfaction', 'int'),
      ('MonthlyIncome', 'int'),
      ('NumCompaniesWorked', 'int'),
      ('PercentSalaryHike', 'int'),
      ('PerformanceRating', 'int'),
      ('TotalWorkingYears', 'int'),
      ('TrainingTimesLastYear', 'int'),
      ('WorkLifeBalance', 'int'),
      ('YearsAtCompany', 'int'),
      ('AttritionMod', 'int'),
      ('MaritalStatusMod', 'int'),
      ('GenderMod', 'int')]
```

```
[33]: dffinal1=df_final.toPandas()
```

```
[36]: dffinal1
```

```
[36]:
```

	Age	DistanceFromHome	EnvironmentSatisfaction	JobInvolvement	\
0	41	1	2	3	
1	49	8	3	2	
2	37	2	4	2	
3	33	3	4	3	
4	27	2	1	3	
...	

1465	36	23	3	4
1466	39	6	4	2
1467	27	4	2	4
1468	49	2	4	2
1469	34	8	2	4

	JobLevel	JobSatisfaction	MonthlyIncome	NumCompaniesWorked	\
0	2	4	5993	8	
1	2	2	5130	1	
2	1	3	2090	6	
3	1	3	2909	1	
4	1	2	3468	9	
...	
1465	2	4	2571	4	
1466	3	1	9991	4	
1467	2	2	6142	1	
1468	2	2	5390	2	
1469	2	3	4404	2	

	PercentSalaryHike	PerformanceRating	TotalWorkingYears	\
0	11	3	8	
1	23	4	10	
2	15	3	7	
3	11	3	8	
4	12	3	6	
...	
1465	17	3	17	
1466	15	3	9	
1467	20	4	6	
1468	14	3	17	
1469	12	3	6	

	TrainingTimesLastYear	WorkLifeBalance	YearsAtCompany	AttritionMod	\
0	0	1	6	1	
1	3	3	10	0	
2	3	3	0	1	
3	3	3	8	0	
4	3	3	2	0	
...	
1465	3	3	5	0	
1466	5	3	7	0	
1467	0	3	6	0	
1468	3	2	9	0	
1469	3	4	4	0	

	MaritalStatusMod	GenderMod
0	0	0

```

1           0           1
2           0           1
3           0           0
4           0           1
...         ...         ...
1465        0           1
1466        0           1
1467        0           1
1468        0           1
1469        0           1

```

[1470 rows x 17 columns]

```

[37]: featurescols =
    ↳ ['Age', 'DistanceFromHome', 'EnvironmentSatisfaction', 'JobInvolvement', 'JobLevel',
        ↳
    ↳ 'JobSatisfaction', 'MonthlyIncome', 'NumCompaniesWorked', 'PercentSalaryHike', 'PerformanceRati
        ↳
    ↳ 'TrainingTimesLastYear', 'WorkLifeBalance', 'YearsAtCompany', 'MaritalStatusMod', 'GenderMod']
from pyspark.ml.feature import VectorAssembler, StringIndexer, VectorIndexer,
    ↳ MinMaxScaler, Normalizer

from pyspark.ml.classification import LogisticRegression

Vect =
    ↳ VectorAssembler(inputCols=['Age', 'DistanceFromHome', 'EnvironmentSatisfaction', 'JobInvolvement
        ↳
    ↳ 'JobSatisfaction', 'MonthlyIncome', 'NumCompaniesWorked', 'PercentSalaryHike', 'PerformanceRati
        ↳
    ↳ 'TrainingTimesLastYear', 'WorkLifeBalance', 'YearsAtCompany', 'MaritalStatusMod', 'GenderMod'],
    ↳ outputCol="features")

```

```

[38]: catIdx = VectorIndexer(inputCol = Vect.getOutputCol(), outputCol =
    ↳ "idxCatFeatures")

```

```

[39]: train, test = df_final.randomSplit([0.8,0.2])

```

```

[40]: train.count()

```

[40]: 1177

```

[41]: test.count()

```

[41]: 293

0	0.0	
1	0.0	
1	0.0	
0	0.0	
1	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
1	0.0	
1	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
1	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
1	0.0	
1	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
1	0.0	
1	0.0	
0	0.0	
1	0.0	
0	0.0	
1	0.0	
0	0.0	

0	0.0	
1	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
1	0.0	
0	0.0	
1	0.0	
0	0.0	
0	0.0	
0	0.0	
0	0.0	
1	0.0	
0	0.0	
0	0.0	
0	0.0	
1	0.0	
0	0.0	
0	0.0	
0	0.0	
1	0.0	
0	0.0	

+-----+-----+

only showing top 100 rows

```
[114]: tp = float(predicted.filter("prediction == 1.0 AND AttritionMod== 1").count())
fp = float(predicted.filter("prediction == 1.0 AND AttritionMod == 0").count())
tn = float(predicted.filter("prediction == 0.0 AND AttritionMod== 0").count())
fn = float(predicted.filter("prediction == 0.0 AND AttritionMod == 1").count())
pr = tp / (tp + fp)
re = tp / (tp + fn)
metrics = spark.createDataFrame([
    ("TP", tp),
    ("FP", fp),
    ("TN", tn),
    ("FN", fn),
    ("Precision", pr),
    ("Recall", re),
    ("F1", 2*pr*re/(re+pr))],["metric", "value"])
metrics.show()
```

```

ZeroDivisionError                                Traceback (most recent call last)
/tmp/ipykernel_8107/768638879.py in <module>
      3 tn = float(predicted.filter("prediction == 0.0 AND AttritionMod== 0").
↪count())
      4 fn = float(predicted.filter("prediction == 0.0 AND AttritionMod == 1").
↪count())
----> 5 pr = tp / (tp + fp)
      6 re = tp / (tp + fn)
      7 metrics = spark.createDataFrame([

ZeroDivisionError: float division by zero

```

```

[ ]: evaluator = BinaryClassificationEvaluator(labelCol="AttritionMod",
↪rawPredictionCol="rawPrediction", metricName="areaUnderROC")
aur = evaluator.evaluate(prediction)
print ("AUR = ", aur)

```

```

[61]: from pyspark.ml.feature import VectorAssembler
assembler = VectorAssembler(inputCols=featureCols, outputCol="features")

assembled_df = assembler.transform(df_final)
assembled_df.show(10, truncate=False)

```

```

+---+-----+-----+-----+-----+-----+-----+
|Age|DistanceFromHome|EnvironmentSatisfaction|JobInvolvement|JobLevel|JobSatisfac|
tion|MonthlyIncome|NumCompaniesWorked|PercentSalaryHike|PerformanceRating|Total
WorkingYears|TrainingTimesLastYear|WorkLifeBalance|YearsAtCompany|AttritionMod|M
aritalStatusMod|GenderMod|features
|
+---+-----+-----+-----+-----+-----+-----+
|41|1|2|3|2|4|
|5993|8|11|3|2|8|
|0|1|6|1|0|
|0|
|[41.0,1.0,2.0,3.0,2.0,4.0,5993.0,8.0,11.0,3.0,8.0,0.0,1.0,6.0,0.0,0.0]|
|49|8|3|2|2|2|
|5130|1|23|4|10|
|3|3|10|0|0|
|1|

```

```

| [49.0,8.0,3.0,2.0,2.0,2.0,5130.0,1.0,23.0,4.0,10.0,3.0,3.0,10.0,0.0,1.0] |
| 37 | 2 | | | 4 | | | 2 | | | 1 | | | 3 |
| 2090 | | 6 | | | 15 | | | 3 | | | 7 |
| 3 | | | 3 | | | 0 | | | 1 | | | 0 |
| 1 |
| [37.0,2.0,4.0,2.0,1.0,3.0,2090.0,6.0,15.0,3.0,7.0,3.0,3.0,0.0,0.0,1.0] |
| 33 | 3 | | | 4 | | | 3 | | | 1 | | | 3 |
| 2909 | | 1 | | | 11 | | | 3 | | | 8 |
| 3 | | | 3 | | | 8 | | | 0 | | | 0 |
| 0 |
| [33.0,3.0,4.0,3.0,1.0,3.0,2909.0,1.0,11.0,3.0,8.0,3.0,3.0,8.0,0.0,0.0] |
| 27 | 2 | | | 1 | | | 3 | | | 1 | | | 2 |
| 3468 | | 9 | | | 12 | | | 3 | | | 6 |
| 3 | | | 3 | | | 2 | | | 0 | | | 0 |
| 1 |
| [27.0,2.0,1.0,3.0,1.0,2.0,3468.0,9.0,12.0,3.0,6.0,3.0,3.0,2.0,0.0,1.0] |
| 32 | 2 | | | 4 | | | 3 | | | 1 | | | 4 |
| 3068 | | 0 | | | 13 | | | 3 | | | 8 |
| 2 | | | 2 | | | 7 | | | 0 | | | 0 |
| 1 |
| [32.0,2.0,4.0,3.0,1.0,4.0,3068.0,0.0,13.0,3.0,8.0,2.0,2.0,7.0,0.0,1.0] |
| 59 | 3 | | | 3 | | | 4 | | | 1 | | | 1 |
| 2670 | | 4 | | | 20 | | | 4 | | | 12 |
| 3 | | | 2 | | | 1 | | | 0 | | | 0 |
| 0 |
| [59.0,3.0,3.0,4.0,1.0,1.0,2670.0,4.0,20.0,4.0,12.0,3.0,2.0,1.0,0.0,0.0] |
| 30 | 24 | | | 4 | | | 3 | | | 1 | | | 3 |
| 2693 | | 1 | | | 22 | | | 4 | | | 1 |
| 2 | | | 3 | | | 1 | | | 0 | | | 0 |
| 1 |
| [30.0,24.0,4.0,3.0,1.0,3.0,2693.0,1.0,22.0,4.0,1.0,2.0,3.0,1.0,0.0,1.0] |
| 38 | 23 | | | 4 | | | 2 | | | 3 | | | 3 |
| 9526 | | 0 | | | 21 | | | 4 | | | 10 |
| 2 | | | 3 | | | 9 | | | 0 | | | 0 |
| 1 |
| [38.0,23.0,4.0,2.0,3.0,3.0,9526.0,0.0,21.0,4.0,10.0,2.0,3.0,9.0,0.0,1.0] |
| 36 | 27 | | | 3 | | | 3 | | | 2 | | | 3 |
| 5237 | | 6 | | | 13 | | | 3 | | | 17 |
| 3 | | | 2 | | | 7 | | | 0 | | | 0 |
| 1 |
| [36.0,27.0,3.0,3.0,2.0,3.0,5237.0,6.0,13.0,3.0,17.0,3.0,2.0,7.0,0.0,1.0] |
+---+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+

```

only showing top 10 rows

```
[62]: normalizer = Normalizer(inputCol="features",outputCol="features_norm")

[63]: lr = LogisticRegression(featuresCol="features_norm",labelCol="AttritionMod",maxIter=10,regParam=
↳3,elasticNetParam=0.2)

[64]: pipeline = Pipeline(stages=[assembler,normalizer,lr])

[65]: pipelineModel = pipeline.fit(train)

[66]: prediction = pipelineModel.transform(test)

[67]: prediction.select("AttritionMod","prediction").show(100)
```

```
+-----+-----+
|AttritionMod|prediction|
+-----+-----+
|          0|        0.0|
|          1|        0.0|
|          1|        0.0|
|          1|        0.0|
|          0|        0.0|
|          0|        0.0|
|          0|        0.0|
|          0|        0.0|
|          1|        0.0|
|          1|        0.0|
|          0|        0.0|
|          1|        0.0|
|          1|        0.0|
|          0|        0.0|
|          0|        0.0|
|          1|        0.0|
|          1|        0.0|
|          0|        0.0|
|          0|        0.0|
|          1|        0.0|
|          0|        0.0|
|          0|        0.0|
|          0|        0.0|
|          0|        0.0|
|          0|        0.0|
|          1|        0.0|
|          0|        0.0|
|          1|        0.0|
```

[illegible]

	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	1	0.0
	1	0.0
	1	0.0
	0	0.0
	0	0.0
	1	0.0
	0	0.0
	0	0.0
	1	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0

```

+-----+-----+
only showing top 100 rows

```

```

[68]: from pyspark.ml.feature import VectorAssembler
assembler = VectorAssembler(inputCols=featureCols, outputCol="features")

assembled_df = assembler.transform(df_final)
assembled_df.show(10, truncate=False)

```

```

+---+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+
|Age|DistanceFromHome|EnvironmentSatisfaction|JobInvolvement|JobLevel|JobSatisfac
tion|MonthlyIncome|NumCompaniesWorked|PercentSalaryHike|PerformanceRating|Total
WorkingYears|TrainingTimesLastYear|WorkLifeBalance|YearsAtCompany|AttritionMod|M
aritalStatusMod|GenderMod|features
|
+---+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+
|41 |1|                                     |3|                                     |2|                                     |4|                                     |3|

```


5993	8	11	3	8
0	1	6	1	0
0				
[41.0,1.0,2.0,3.0,2.0,4.0,5993.0,8.0,11.0,3.0,8.0,0.0,1.0,6.0,0.0,0.0]				
49	8	3	2	2
5130	1	23	4	10
3	3	10	0	0
1				
[49.0,8.0,3.0,2.0,2.0,2.0,5130.0,1.0,23.0,4.0,10.0,3.0,3.0,10.0,0.0,1.0]				
37	2	4	2	1
2090	6	15	3	7
3	3	0	1	0
1				
[37.0,2.0,4.0,2.0,1.0,3.0,2090.0,6.0,15.0,3.0,7.0,3.0,3.0,0.0,0.0,1.0]				
33	3	4	3	1
2909	1	11	3	8
3	3	8	0	0
0				
[33.0,3.0,4.0,3.0,1.0,3.0,2909.0,1.0,11.0,3.0,8.0,3.0,3.0,8.0,0.0,0.0]				
27	2	1	3	1
3468	9	12	3	6
3	3	2	0	0
1				
[27.0,2.0,1.0,3.0,1.0,2.0,3468.0,9.0,12.0,3.0,6.0,3.0,3.0,2.0,0.0,1.0]				
32	2	4	3	1
3068	0	13	3	8
2	2	7	0	0
1				
[32.0,2.0,4.0,3.0,1.0,4.0,3068.0,0.0,13.0,3.0,8.0,2.0,2.0,7.0,0.0,1.0]				
59	3	3	4	1
2670	4	20	4	12
3	2	1	0	0
0				
[59.0,3.0,3.0,4.0,1.0,1.0,2670.0,4.0,20.0,4.0,12.0,3.0,2.0,1.0,0.0,0.0]				
30	24	4	3	1
2693	1	22	4	1
2	3	1	0	0
1				
[30.0,24.0,4.0,3.0,1.0,3.0,2693.0,1.0,22.0,4.0,1.0,2.0,3.0,1.0,0.0,1.0]				
38	23	4	2	3
9526	0	21	4	10
2	3	9	0	0
1				
[38.0,23.0,4.0,2.0,3.0,3.0,9526.0,0.0,21.0,4.0,10.0,2.0,3.0,9.0,0.0,1.0]				
36	27	3	3	2
5237	6	13	3	17
3	2	7	0	0
1				

```
| [36.0,27.0,3.0,3.0,2.0,3.0,5237.0,6.0,13.0,3.0,17.0,3.0,2.0,7.0,0.0,1.0] |
+---+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 10 rows
```

```
[69]: standardScaler = StandardScaler(inputCol="features",
    ↪outputCol="features_scaled")
scaled_df = standardScaler.fit(assembled_df).transform(assembled_df)

scaled_df.select("features", "features_scaled").show(10, truncate=False)
train_data, test_data = scaled_df.randomSplit([0.8,0.2])
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
|features
|features_scaled
|
+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+
| [41.0,1.0,2.0,3.0,2.0,4.0,5993.0,8.0,11.0,3.0,8.0,0.0,1.0,6.0,0.0,0.0] | [4.488
048578282528,0.12335225387518614,1.8296885387233528,4.216081821875181,1.80678282
70748993,3.6269792461350723,1.2729513621525657,3.2025505034441926,3.005515626752
296,8.31431377233402,1.0281743317865935,0.0,1.4154765923196735,0.979347974707337
7,0.0,0.0]
|
| [49.0,8.0,3.0,2.0,2.0,2.0,5130.0,1.0,23.0,4.0,10.0,3.0,3.0,10.0,0.0,1.0] | [5.363
765374044972,0.9868180310014891,2.744532808085029,2.810721214583454,1.8067828270
748993,1.8134896230675361,1.089644666751654,0.40031881293052407,6.28425994684571
,11.085751696445358,1.2852179147332419,2.3268970467566765,4.246429776959021,1.63
22466245122293,0.0,2.0405470343872216]
|
| [37.0,2.0,4.0,2.0,1.0,3.0,2090.0,6.0,15.0,3.0,7.0,3.0,3.0,0.0,0.0,1.0] | [4.050
190180401305,0.24670450775037228,3.6593770774467056,2.810721214583454,0.90339141
35374497,2.7202344346013043,0.4439293086765997,2.401912877583144,4.0984304001167
68,8.31431377233402,0.8996525403132692,2.3268970467566765,4.246429776959021,0.0,
0.0,2.0405470343872216]
|
| [33.0,3.0,4.0,3.0,1.0,3.0,2909.0,1.0,11.0,3.0,8.0,3.0,3.0,8.0,0.0,0.0] | [3.612
331782520083,0.37005676162555845,3.6593770774467056,4.216081821875181,0.90339141
35374497,2.7202344346013043,0.6178901238948462,0.40031881293052407,3.00551562675
2296,8.31431377233402,1.0281743317865935,2.3268970467566765,4.246429776959021,1.
```

```

3057972996097835,0.0,0.0] |
| [27.0,2.0,1.0,3.0,1.0,2.0,3468.0,9.0,12.0,3.0,6.0,3.0,3.0,2.0,0.0,1.0] | [2.955
54418569825,0.24670450775037228,0.9148442693616764,4.216081821875181,0.903391413
5374497,1.8134896230675361,0.7366252834882525,3.6028693163747167,3.2787443200934
137,8.31431377233402,0.771130748839945,2.3268970467566765,4.246429776959021,0.32
64493249024459,0.0,2.0405470343872216] |
| [32.0,2.0,4.0,3.0,1.0,4.0,3068.0,0.0,13.0,3.0,8.0,2.0,2.0,7.0,0.0,1.0] | [3.502
8671830497777,0.24670450775037228,3.6593770774467056,4.216081821875181,0.9033914
135374497,3.6269792461350723,0.6516627363731139,0.0,3.5519730134345315,8.3143137
7233402,1.0281743317865935,1.5512646978377842,2.830953184639347,1.14257263715856
05,0.0,2.0405470343872216] |
| [59.0,3.0,3.0,4.0,1.0,1.0,2670.0,4.0,20.0,4.0,12.0,3.0,2.0,1.0,0.0,0.0] | [6.458
411368748028,0.37005676162555845,2.744532808085029,5.621442429166908,0.903391413
5374497,0.9067448115337681,0.5671250019935509,1.6012752517220963,5.4645738668223
57,11.085751696445358,1.54226149767989,2.3268970467566765,2.830953184639347,0.16
322466245122294,0.0,0.0] |
| [30.0,24.0,4.0,3.0,1.0,3.0,2693.0,1.0,22.0,4.0,1.0,2.0,3.0,1.0,0.0,1.0] | [3.283
9379841091665,2.9604540930044676,3.6593770774467056,4.216081821875181,0.90339141
35374497,2.7202344346013043,0.5720103484526713,0.40031881293052407,6.01103125350
4592,11.085751696445358,0.12852179147332418,1.5512646978377842,4.246429776959021
,0.16322466245122294,0.0,2.0405470343872216] |
| [38.0,23.0,4.0,2.0,3.0,3.0,9526.0,0.0,21.0,4.0,10.0,2.0,3.0,9.0,0.0,1.0] | [4.159
6547798716115,2.8371018391292813,3.6593770774467056,2.810721214583454,2.71017424
0612349,2.7202344346013043,2.023383059547028,0.0,5.737802560163474,11.0857516964
45358,1.2852179147332419,1.5512646978377842,4.246429776959021,1.4690219620610065
,0.0,2.0405470343872216] |
| [36.0,27.0,3.0,3.0,2.0,3.0,5237.0,6.0,13.0,3.0,17.0,3.0,2.0,7.0,0.0,1.0] | [3.940
725580931,3.3305108546300257,2.744532808085029,4.216081821875181,1.8067828270748
993,2.7202344346013043,1.1123721481049536,2.401912877583144,3.5519730134345315,8
.31431377233402,2.1848704550465112,2.3268970467566765,2.830953184639347,1.142572
6371585605,0.0,2.0405470343872216] |
+-----+
-----
-----
-----
-----+
only showing top 10 rows

```

```

[70]: lr = LogisticRegression(featuresCol="features_scaled",labelCol="AttritionMod",maxIter=10,regParam=0.2,elasticNetParam=0.2)

linearModel = lr.fit(train_data)
predictions = linearModel.transform(test_data)

predictions.select("AttritionMod","prediction").show(100)

```

+-----+		+-----+
AttritionMod		prediction
+-----+		+-----+
	1	0.0
	1	0.0
	0	0.0
	0	0.0
	0	0.0
	1	0.0
	0	0.0
	1	0.0
	0	0.0
	0	0.0
	0	0.0
	1	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	1	0.0
	0	0.0
	1	0.0
	0	0.0
	1	0.0
	0	0.0
	0	0.0
	1	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	1	0.0
	1	0.0
	0	0.0
	0	0.0
	0	0.0

[illegible]

	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0
	0	0.0

+-----+-----+

only showing top 100 rows

[]: