

KP02-Copy3

November 16, 2021

```
[2]: !pip install pyspark
```

```
Requirement already satisfied: pyspark in /usr/lib/spark/python (3.1.2)
Requirement already satisfied: py4j==0.10.9 in
/opt/conda/miniconda3/lib/python3.8/site-packages (from pyspark) (0.10.9)
WARNING: Running pip as the 'root' user can result in broken permissions
and conflicting behaviour with the system package manager. It is recommended to
use a virtual environment instead: https://pip.pypa.io/warnings/venv
```

```
[3]: import os
import pandas as pd
import numpy as np

from pyspark import SparkConf, SparkContext
from pyspark.sql import SparkSession, SQLContext

import pyspark.sql.functions as F
from pyspark.sql.functions import udf, col

from pyspark.ml.tuning import ParamGridBuilder, CrossValidator
from pyspark.ml.evaluation import BinaryClassificationEvaluator

from pyspark.sql.functions import udf, col

from pyspark.mllib.evaluation import RegressionMetrics

from pyspark.ml.tuning import ParamGridBuilder, CrossValidator,
↳CrossValidatorModel
from pyspark.ml.feature import VectorAssembler, StandardScaler
from pyspark.ml.evaluation import RegressionEvaluator

from pyspark.ml.tuning import ParamGridBuilder, CrossValidator
from pyspark.ml.evaluation import BinaryClassificationEvaluator

from pyspark.ml import Pipeline
```

```
[4]: from pyspark.sql import SparkSession
      spark = SparkSession \
            .builder \
            .getOrCreate()
```

```
[5]: df = spark.read.format('csv').option("header", "true").option("delimiter", ",").
      ↪load("gs://bdamlproject/TurnoverIntentionatBenzInfotech.csv")
```

```
[6]: for col in df.columns:
      print(col)
```

```
Tenure
PerfRating
JobSatisfaction
Engagement
OnsiteOpportunity
Awards
Flexihours
Intention
```

```
[6]: df.count()
```

```
[6]: 361
```

```
[7]: df.dtypes
```

```
[7]: [('Tenure', 'string'),
      ('PerfRating', 'string'),
      ('JobSatisfaction', 'string'),
      ('Engagement', 'string'),
      ('OnsiteOpportunity', 'string'),
      ('Awards', 'string'),
      ('Flexihours', 'string'),
      ('Intention', 'string')]
```

```
[8]: df1=df.toPandas()
```

```
[9]: df1
```

```
[9]:   Tenure  PerfRating  JobSatisfaction  Engagement  OnsiteOpportunity  Awards  \
0        2          4          4          1          0          1
1        1          1          3          1          0          1
2        1          1          3          2          1          0
3        1          1          2          4          0          1
4        1          1          3          1          0          1
```

```

..      ...      ...      ...      ...      ...      ...      ...
356      4      3      3      3      0      1
357      2      2      2      3      1      0
358      4      3      2      1      1      0
359      5      1      5      5      1      0
360      2      1      1      1      1      1

```

```

      Flexihours Intention
0          0          1
1          1          1
2          1          1
3          1          1
4          1          1
..      ...      ...
356      1          1
357      1          1
358      1          0
359      0          0
360      1          1

```

[361 rows x 8 columns]

```
[10]: df1.describe()
```

```

[10]:      Tenure PerfRating JobSatisfaction Engagement OnsiteOpportunity Awards \
count      361      361      361      361      361      361
unique       8       5       5       5       2       2
top          2       1       4       4       1       0
freq        61      90      85      86     194     182

```

```

      Flexihours Intention
count      361      361
unique       2       2
top          0       0
freq       196     181

```

```
[11]: df1.isna().sum()
```

```

[11]: Tenure          0
      PerfRating      0
      JobSatisfaction  0
      Engagement      0
      OnsiteOpportunity 0
      Awards          0
      Flexihours       0
      Intention        0
      dtype: int64

```

```
[1]: df= df.withColumn("Tenure",df["Tenure"].cast(IntegerType()))
df= df.withColumn("PerfRating",df["PerfRating"].cast(IntegerType()))
df= df.withColumn("JobSatisfaction",df["JobSatisfaction"].cast(IntegerType()))
df= df.withColumn("Engagement",df["Engagement"].cast(IntegerType()))
df= df.withColumn("OnsiteOpportunity",df["OnsiteOpportunity"].
    ↳cast(IntegerType()))
df= df.withColumn("Awards",df["Awards"].cast(IntegerType()))
df= df.withColumn("Flexihours",df["Flexihours"].cast(IntegerType()))
df= df.withColumn("Intention",df["Intention"].cast(IntegerType()))
```

```
-----
NameError                                Traceback (most recent call last)
/tmp/ipykernel_8947/2029380153.py in <module>
----> 1 df= df.withColumn("Tenure",df["Tenure"].cast(IntegerType()))
      2 df= df.withColumn("PerfRating",df["PerfRating"].cast(IntegerType()))
      3 df= df.withColumn("JobSatisfaction",df["JobSatisfaction"].
    ↳cast(IntegerType()))
      4 df= df.withColumn("Engagement",df["Engagement"].cast(IntegerType()))
      5 df= df.withColumn("OnsiteOpportunity",df["OnsiteOpportunity"].
    ↳cast(IntegerType()))

NameError: name 'df' is not defined
```

```
[ ]: from pyspark.ml.stat import Correlation
from pyspark.ml.feature import VectorAssembler
```

```
[43]: from pyspark.sql.types import IntegerType
```

```
[44]: from pyspark.sql.functions import when
from pyspark.sql.functions import lit
```

```
[45]: df.columns
```

```
[45]: ['Tenure',
      'PerfRating',
      'JobSatisfaction',
      'Engagement',
      'OnsiteOpportunity',
      'Awards',
      'Flexihours',
      'Intention']
```

```
[46]: features = ['Tenure',
      'PerfRating',
      'JobSatisfaction',
      'Engagement',
```

```

'OnsiteOpportunity',
'Awards',
'Flexihours']
from pyspark.ml.feature import VectorAssembler, StringIndexer, VectorIndexer,
↳MinMaxScaler,Normalizer

from pyspark.ml.classification import LogisticRegression

Vect = VectorAssembler(inputCols=['Tenure',
'PerfRating',
'JobSatisfaction',
'Engagement',
'OnsiteOpportunity',
'Awards',
'Flexihours'], outputCol="features")

```

```

[47]: catIdx = VectorIndexer(inputCol = Vect.getOutputCol(), outputCol =
↳"idxCatFeatures")

```

```

[48]: train, test = df.randomSplit([0.8,0.2])

```

```

[49]: train.count()

```

```

[49]: 278

```

```

[50]: test.count()

```

```

[50]: 83

```

```

[51]: lr =
↳LogisticRegression(labelCol='Intention',featuresCol="features",maxIter=10,regParam=0.
↳3)

```

```

[52]: pipeline1 = Pipeline(stages=[Vect, catIdx, lr])

```

```

[53]: pipelineModel = pipeline1.fit(train)

```

```

21/11/13 01:19:35 WARN com.github.fommil.netlib.BLAS: Failed to load
implementation from: com.github.fommil.netlib.NativeSystemBLAS
21/11/13 01:19:35 WARN com.github.fommil.netlib.BLAS: Failed to load
implementation from: com.github.fommil.netlib.NativeRefBLAS

```

```

[54]: prediction1 = pipelineModel.transform(test)

```

```

[56]: predicted=prediction1.select("Intention","prediction")
predicted.show(100, truncate=False)

```

+-----+-----+		
Intention prediction		
+-----+-----+		
1	1.0	
1	1.0	
1	1.0	
0	1.0	
1	1.0	
0	0.0	
1	1.0	
0	1.0	
0	1.0	
1	1.0	
1	1.0	
1	1.0	
1	1.0	
1	1.0	
1	1.0	
1	1.0	
1	1.0	
1	1.0	
1	1.0	
0	1.0	
1	0.0	
1	0.0	
0	0.0	
1	0.0	
1	1.0	
1	1.0	
0	1.0	
1	1.0	
0	1.0	
1	1.0	
1	1.0	
1	1.0	
1	1.0	
1	1.0	
1	1.0	
1	0.0	
1	1.0	
1	1.0	
1	1.0	
0	1.0	
1	1.0	
1	1.0	
0	0.0	
0	0.0	
1	1.0	
0	0.0	
0	0.0	

$$+ \text{---} + \text{---} +$$

```
tp = float(predicted.filter("prediction == 1.0 AND Intention== 1").count())
fp = float(predicted.filter("prediction == 1.0 AND Intention == 0").count())
tn = float(predicted.filter("prediction == 0.0 AND Intention== 0").count())
fn = float(predicted.filter("prediction == 0.0 AND Intention == 1").count())
pr = tp / (tp + fp)
re = tp / (tp + fn)
```

```
metrics = spark.createDataFrame([
    ("TP", tp),
    ("FP", fp),
    ("TN", tn),
    ("FN", fn),
    ("Precision", pr),
    ("Recall", re),
    ("F1", 2*pr*re/(re+pr))],["metric", "value"])
metrics.show()
```

[Stage 46:>

(0 + 1) / 1]

```
+-----+-----+
| metric|      value|
+-----+-----+
|      TP|      34.0|
|      FP|      10.0|
|      TN|      30.0|
|      FN|       9.0|
|Precision|0.77272727272727|
|  Recall|0.7906976744186046|
|       F1|0.7816091954022988|
+-----+-----+
```

```
[59]: evaluator = BinaryClassificationEvaluator(labelCol="AttritionMod",
    ↳rawPredictionCol="rawPrediction", metricName="areaUnderROC")
aur = evaluator.evaluate(prediction)
print ("AUR = ", aur)
```

```
-----
NameError                                Traceback (most recent call last)
/tmp/ipykernel_5439/3420554418.py in <module>
      1 evaluator = BinaryClassificationEvaluator(labelCol="AttritionMod",
    ↳rawPredictionCol="rawPrediction", metricName="areaUnderROC")
----> 2 aur = evaluator.evaluate(prediction)
      3 print ("AUR = ", aur)

NameError: name 'prediction' is not defined
```

```
[61]: from pyspark.ml.feature import VectorAssembler
assembler = VectorAssembler(inputCols=featureCols, outputCol="features")

assembled_df = assembler.transform(df_final)
assembled_df.show(10, truncate=False)
```



```

+---+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+
-----+
|Age|DistanceFromHome|EnvironmentSatisfaction|JobInvolvement|JobLevel|JobSatisfac
tion|MonthlyIncome|NumCompaniesWorked|PercentSalaryHike|PerformanceRating|Total
WorkingYears|TrainingTimesLastYear|WorkLifeBalance|YearsAtCompany|AttritionMod|M
aritalStatusMod|GenderMod|features
|
+---+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+
-----+
|41 |1          |2          |3          |2          |4
|5993          |8          |11         |3          |8
|0            |1          |6          |1          |0
|0
|[41.0,1.0,2.0,3.0,2.0,4.0,5993.0,8.0,11.0,3.0,8.0,0.0,1.0,6.0,0.0,0.0] |
|49 |8          |3          |2          |2          |2
|5130          |1          |23         |4          |10
|3            |3          |10         |0          |0
|1
|[49.0,8.0,3.0,2.0,2.0,2.0,5130.0,1.0,23.0,4.0,10.0,3.0,3.0,10.0,0.0,1.0] |
|37 |2          |4          |2          |1          |3
|2090          |6          |15         |3          |7
|3            |3          |0          |1          |0
|1
|[37.0,2.0,4.0,2.0,1.0,3.0,2090.0,6.0,15.0,3.0,7.0,3.0,3.0,0.0,0.0,1.0] |
|33 |3          |4          |3          |1          |3
|2909          |1          |11         |3          |8
|3            |3          |8          |0          |0
|0
|[33.0,3.0,4.0,3.0,1.0,3.0,2909.0,1.0,11.0,3.0,8.0,3.0,3.0,8.0,0.0,0.0] |
|27 |2          |1          |3          |1          |2
|3468          |9          |12         |3          |6
|3            |3          |2          |0          |0
|1
|[27.0,2.0,1.0,3.0,1.0,2.0,3468.0,9.0,12.0,3.0,6.0,3.0,3.0,2.0,0.0,1.0] |
|32 |2          |4          |3          |1          |4
|3068          |0          |13         |3          |8
|2            |2          |7          |0          |0
|1
|[32.0,2.0,4.0,3.0,1.0,4.0,3068.0,0.0,13.0,3.0,8.0,2.0,2.0,7.0,0.0,1.0] |
|59 |3          |3          |4          |1          |1
|2670          |4          |20         |4          |12
|3            |2          |1          |0          |0

```

```

|0
|[59.0,3.0,3.0,4.0,1.0,1.0,2670.0,4.0,20.0,4.0,12.0,3.0,2.0,1.0,0.0,0.0] |
|30 |24          |4          |3          |1          |3
|2693          |1          |22          |4          |1
|2          |3          |1          |0          |0
|1
|[30.0,24.0,4.0,3.0,1.0,3.0,2693.0,1.0,22.0,4.0,1.0,2.0,3.0,1.0,0.0,1.0] |
|38 |23          |4          |2          |3          |3
|9526          |0          |21          |4          |10
|2          |3          |9          |0          |0
|1
|[38.0,23.0,4.0,2.0,3.0,3.0,9526.0,0.0,21.0,4.0,10.0,2.0,3.0,9.0,0.0,1.0] |
|36 |27          |3          |3          |2          |3
|5237          |6          |13          |3          |17
|3          |2          |7          |0          |0
|1
|[36.0,27.0,3.0,3.0,2.0,3.0,5237.0,6.0,13.0,3.0,17.0,3.0,2.0,7.0,0.0,1.0] |
+---+-----+-----+-----+-----+-----+-----+-----+
+---+-----+-----+-----+-----+-----+-----+-----+
+---+-----+-----+-----+-----+-----+-----+-----+
+---+-----+-----+-----+-----+-----+-----+-----+
+---+-----+
only showing top 10 rows

```

```
[62]: normalizer = Normalizer(inputCol="features",outputCol="features_norm")
```

```
[63]: lr =
↳ LogisticRegression(featuresCol="features_norm",labelCol="AttritionMod",maxIter=10,regParam=
↳ 3,elasticNetParam=0.2)
```

```
[64]: pipeline = Pipeline(stages=[assembler,normalizer,lr])
```

```
[65]: pipelineModel = pipeline.fit(train)
```

```
[66]: prediction = pipelineModel.transform(test)
```

```
[67]: prediction.select("AttritionMod","prediction").show(100)
```

```

+-----+-----+
|AttritionMod|prediction|
+-----+-----+
|          0|          0.0|
|          1|          0.0|
|          1|          0.0|
|          1|          0.0|
|          0|          0.0|
|          0|          0.0|

```

[illegible]

[illegible]

only showing top 100 rows

```
[68]: from pyspark.ml.feature import VectorAssembler
assembler = VectorAssembler(inputCols=featureCols, outputCol="features")

assembled_df = assembler.transform(df_final)
assembled_df.show(10, truncate=False)
```

```
+---+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+
|Age|DistanceFromHome|EnvironmentSatisfaction|JobInvolvement|JobLevel|JobSatisfac
tion|MonthlyIncome|NumCompaniesWorked|PercentSalaryHike|PerformanceRating|Total
WorkingYears|TrainingTimesLastYear|WorkLifeBalance|YearsAtCompany|AttritionMod|M
aritalStatusMod|GenderMod|features
|
+---+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+
|41 |1          |2          |3          |2          |4
|5993          |8          |11         |3          |8
|0            |1          |6          |1          |0
|0
|[41.0,1.0,2.0,3.0,2.0,4.0,5993.0,8.0,11.0,3.0,8.0,0.0,1.0,6.0,0.0,0.0] |
|49 |8          |3          |2          |2          |2
|5130         |1          |23         |4          |10
|3            |3          |10         |0          |0
|1
|[49.0,8.0,3.0,2.0,2.0,2.0,5130.0,1.0,23.0,4.0,10.0,3.0,3.0,10.0,0.0,1.0] |
|37 |2          |4          |2          |1          |3
|2090         |6          |15         |3          |7
|3            |3          |0          |1          |0
|1
|[37.0,2.0,4.0,2.0,1.0,3.0,2090.0,6.0,15.0,3.0,7.0,3.0,3.0,0.0,0.0,1.0] |
|33 |3          |4          |3          |1          |3
|2909         |1          |11         |3          |8
|3            |3          |8          |0          |0
|0
|[33.0,3.0,4.0,3.0,1.0,3.0,2909.0,1.0,11.0,3.0,8.0,3.0,3.0,8.0,0.0,0.0] |
|27 |2          |1          |3          |1          |2
|3468         |9          |12         |3          |6
|3            |3          |2          |0          |0
|1
|[27.0,2.0,1.0,3.0,1.0,2.0,3468.0,9.0,12.0,3.0,6.0,3.0,3.0,2.0,0.0,1.0] |
```

```

|32 |2          |4          |3          |1          |4
|3068          |0          |13          |3          |8
|2          |2          |7          |0          |0
|1
|[32.0,2.0,4.0,3.0,1.0,4.0,3068.0,0.0,13.0,3.0,8.0,2.0,2.0,7.0,0.0,1.0] |
|59 |3          |3          |4          |1          |1
|2670          |4          |20          |4          |12
|3          |2          |1          |0          |0
|0
|[59.0,3.0,3.0,4.0,1.0,1.0,2670.0,4.0,20.0,4.0,12.0,3.0,2.0,1.0,0.0,0.0] |
|30 |24          |4          |3          |1          |3
|2693          |1          |22          |4          |1
|2          |3          |1          |0          |0
|1
|[30.0,24.0,4.0,3.0,1.0,3.0,2693.0,1.0,22.0,4.0,1.0,2.0,3.0,1.0,0.0,1.0] |
|38 |23          |4          |2          |3          |3
|9526          |0          |21          |4          |10
|2          |3          |9          |0          |0
|1
|[38.0,23.0,4.0,2.0,3.0,3.0,9526.0,0.0,21.0,4.0,10.0,2.0,3.0,9.0,0.0,1.0] |
|36 |27          |3          |3          |2          |3
|5237          |6          |13          |3          |17
|3          |2          |7          |0          |0
|1
|[36.0,27.0,3.0,3.0,2.0,3.0,5237.0,6.0,13.0,3.0,17.0,3.0,2.0,7.0,0.0,1.0] |
+---+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+
only showing top 10 rows

```

```

[69]: standardScaler = StandardScaler(inputCol="features",
    ↪outputCol="features_scaled")
scaled_df = standardScaler.fit(assembled_df).transform(assembled_df)

scaled_df.select("features", "features_scaled").show(10, truncate=False)
train_data, test_data = scaled_df.randomSplit([0.8,0.2])

```

```

+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+
|features
|features_scaled
|

```

```

+-----+
-----
-----
-----
-----+
| [41.0,1.0,2.0,3.0,2.0,4.0,5993.0,8.0,11.0,3.0,8.0,0.0,1.0,6.0,0.0,0.0] | [4.488
048578282528,0.12335225387518614,1.8296885387233528,4.216081821875181,1.80678282
70748993,3.6269792461350723,1.2729513621525657,3.2025505034441926,3.005515626752
296,8.31431377233402,1.0281743317865935,0.0,1.4154765923196735,0.979347974707337
7,0.0,0.0]
|
| [49.0,8.0,3.0,2.0,2.0,2.0,5130.0,1.0,23.0,4.0,10.0,3.0,3.0,10.0,0.0,1.0] | [5.363
765374044972,0.9868180310014891,2.744532808085029,2.810721214583454,1.8067828270
748993,1.8134896230675361,1.089644666751654,0.40031881293052407,6.28425994684571
,11.085751696445358,1.2852179147332419,2.3268970467566765,4.246429776959021,1.63
22466245122293,0.0,2.0405470343872216]
|
| [37.0,2.0,4.0,2.0,1.0,3.0,2090.0,6.0,15.0,3.0,7.0,3.0,3.0,0.0,0.0,1.0] | [4.050
190180401305,0.24670450775037228,3.6593770774467056,2.810721214583454,0.90339141
35374497,2.7202344346013043,0.4439293086765997,2.401912877583144,4.0984304001167
68,8.31431377233402,0.8996525403132692,2.3268970467566765,4.246429776959021,0.0,
0.0,2.0405470343872216]
|
| [33.0,3.0,4.0,3.0,1.0,3.0,2909.0,1.0,11.0,3.0,8.0,3.0,3.0,8.0,0.0,0.0] | [3.612
331782520083,0.37005676162555845,3.6593770774467056,4.216081821875181,0.90339141
35374497,2.7202344346013043,0.6178901238948462,0.40031881293052407,3.00551562675
2296,8.31431377233402,1.0281743317865935,2.3268970467566765,4.246429776959021,1.
3057972996097835,0.0,0.0]
|
| [27.0,2.0,1.0,3.0,1.0,2.0,3468.0,9.0,12.0,3.0,6.0,3.0,3.0,2.0,0.0,1.0] | [2.955
54418569825,0.24670450775037228,0.9148442693616764,4.216081821875181,0.903391413
5374497,1.8134896230675361,0.7366252834882525,3.6028693163747167,3.2787443200934
137,8.31431377233402,0.771130748839945,2.3268970467566765,4.246429776959021,0.32
64493249024459,0.0,2.0405470343872216]
|
| [32.0,2.0,4.0,3.0,1.0,4.0,3068.0,0.0,13.0,3.0,8.0,2.0,2.0,7.0,0.0,1.0] | [3.502
8671830497777,0.24670450775037228,3.6593770774467056,4.216081821875181,0.9033914
135374497,3.6269792461350723,0.6516627363731139,0.0,3.5519730134345315,8.3143137
7233402,1.0281743317865935,1.5512646978377842,2.830953184639347,1.14257263715856
05,0.0,2.0405470343872216]
|
| [59.0,3.0,3.0,4.0,1.0,1.0,2670.0,4.0,20.0,4.0,12.0,3.0,2.0,1.0,0.0,0.0] | [6.458
411368748028,0.37005676162555845,2.744532808085029,5.621442429166908,0.903391413
5374497,0.9067448115337681,0.5671250019935509,1.6012752517220963,5.4645738668223
57,11.085751696445358,1.54226149767989,2.3268970467566765,2.830953184639347,0.16
322466245122294,0.0,0.0]
|
| [30.0,24.0,4.0,3.0,1.0,3.0,2693.0,1.0,22.0,4.0,1.0,2.0,3.0,1.0,0.0,1.0] | [3.283
9379841091665,2.9604540930044676,3.6593770774467056,4.216081821875181,0.90339141
35374497,2.7202344346013043,0.5720103484526713,0.40031881293052407,6.01103125350
4592,11.085751696445358,0.12852179147332418,1.5512646978377842,4.246429776959021
,0.16322466245122294,0.0,2.0405470343872216]
|
| [38.0,23.0,4.0,2.0,3.0,3.0,9526.0,0.0,21.0,4.0,10.0,2.0,3.0,9.0,0.0,1.0] | [4.159
6547798716115,2.8371018391292813,3.6593770774467056,2.810721214583454,2.71017424
0612349,2.7202344346013043,2.023383059547028,0.0,5.737802560163474,11.0857516964

```

```

45358,1.2852179147332419,1.5512646978377842,4.246429776959021,1.4690219620610065
,0.0,2.0405470343872216]
|
|[36.0,27.0,3.0,3.0,2.0,3.0,5237.0,6.0,13.0,3.0,17.0,3.0,2.0,7.0,0.0,1.0]| [3.940
725580931,3.3305108546300257,2.744532808085029,4.216081821875181,1.8067828270748
993,2.7202344346013043,1.1123721481049536,2.401912877583144,3.5519730134345315,8
.31431377233402,2.1848704550465112,2.3268970467566765,2.830953184639347,1.142572
6371585605,0.0,2.0405470343872216]
|
+-----+-----+
-----
-----
-----
-----+

```

only showing top 10 rows

```

[70]: lr = LogisticRegression(featuresCol="features_scaled",labelCol="AttritionMod",maxIter=10,regParam=
3,elasticNetParam=0.2)

linearModel = lr.fit(train_data)
predictions = linearModel.transform(test_data)

predictions.select("AttritionMod","prediction").show(100)

```

```

+-----+-----+
|AttritionMod|prediction|
+-----+-----+
|          1|         0.0|
|          1|         0.0|
|          0|         0.0|
|          0|         0.0|
|          0|         0.0|
|          1|         0.0|
|          0|         0.0|
|          1|         0.0|
|          0|         0.0|
|          0|         0.0|
|          0|         0.0|
|          1|         0.0|
|          0|         0.0|
|          0|         0.0|
|          0|         0.0|
|          0|         0.0|
|          0|         0.0|
|          0|         0.0|
|          0|         0.0|
|          1|         0.0|

```


[illegible]

-----+

only showing top 100 rows